

Bayesian Model Averaging for Ensemble-Based Estimates of Solvation Free Energies

Luke J. Gosink,[†] Christopher C. Overall,[†] Sarah M. Reehl,[†] Paul D. Whitney,[‡]

David L. Mobley,[¶] and Nathan A. Baker^{*,‡,§}

[†]*Computational and Statistical Analytics Division, Pacific Northwest National Laboratory,
Richland, WA 99352, USA*

[‡]*Advanced Computing, Mathematics, and Data Division, Pacific Northwest National
Laboratory, Richland, WA 99352, USA*

[¶]*Departments of Pharmaceutical Sciences and Chemistry, University of California, Irvine,
Irvine, CA 92697*

[§]*Division of Applied Mathematics, Brown University, Providence, RI 02912, USA*

E-mail: nathan.baker@pnnl.gov

Phone: +1-509-375-3997

Abstract

This paper applies the Bayesian Model Averaging (BMA) statistical ensemble technique to estimate small molecule solvation free energies. There is a wide range of methods available for predicting solvation free energies, ranging from empirical statistical models to *ab initio* quantum mechanical approaches. Each of these methods is based on a set of conceptual assumptions that can affect a method’s predictive accuracy and transferability. Using an iterative statistical process, we have selected and combined solvation energy estimates using an ensemble of 17 diverse methods from the Statistical Assessment of Modeling of Proteins and Ligands (SAMPL) SAMPL4 blind prediction study to form a single, aggregated solvation energy estimate. The ensemble design process evaluates the statistical information in each individual method as well as the performance of the aggregate estimate obtained from the ensemble as a whole. Methods that possess minimal or redundant information are pruned from the ensemble and the evaluation process repeats until aggregate predictive performance can no longer be improved. We show that this process results in a final aggregate estimate that outperforms all individual methods by reducing estimate errors by as much as 91% to 1.2 kcal mol⁻¹ accuracy. We also compare our iterative refinement approach to other statistical ensemble approaches and demonstrate that this iterative process reduces estimate errors by as much as 61%. This work provides a new approach for accurate solvation free energy prediction and lays the foundation for future work on aggregate models that can balance computational cost with predictive accuracy.

1 Introduction

Accurate calculation of solvent-solute interactions is an important component of robust molecular simulation including protein structure prediction,^{1,2,3} conformational ensemble calculations,^{4,5,6} and binding free energy calculations.^{7,8,9} Solvation energy methods for atomically detailed molecular models have a long history of development^{10,11,12,13,14,15,16,17} and have recently benefited from the curation of experimental small molecule solvation data for

blind prediction challenges such as the Statistical Assessment of Modeling of Proteins and Ligands (SAMPL) challenge studies.^{1,2,3,4,5,6,7} By presenting a larger range of solvation free energies and molecular weights than seen in common public datasets, these challenges are helping to advance the development of solvation methods, enabling them to become more robust and accurate in their estimates for a variety of molecular targets.^{8,9,10}

Consistent across many of these challenge studies is the observation that the top-performing methods typically come from a wide range of different strategies.^{11,12,13} Across different challenges, top performers have included explicit solvation methods,^{14,15} implicit solvation methods,^{16,17} and hybrid methods that combine mixed quantum mechanics (QM) with molecular mechanical (MM) approaches.^{18,19} In this context, there is a significant degree of uncertainty associated with how to best select, specify, and evaluate the set of parameters and mathematical systems needed to accurately estimate solvation free energies: e.g., uncertainties due to hydrophobicity,²⁰ surface effects,²¹ and solvent asymmetries.²² This type of *method selection uncertainty* affects a wide range of scientific and mathematical disciplines and is arguably the greatest source of error and risk associated with estimation tasks.^{23,24,25} One of the most powerful ways to address this uncertainty is by combining an ensemble of varied methods (e.g., through a weighted average) to form a single aggregated estimate.^{26,27,28} The motivation behind ensemble approaches is based on two principles: (1) most methods in the ensemble possess some unique, useful information; and (2) no single method is sufficient to account for all uncertainties. When modeled correctly, the information and strengths of individual methods can be combined, and their corresponding weaknesses and biases can be overcome by the strength of the group.^{29,30,31} Ensemble-based estimates are therefore expected to be more reliable and accurate than individual methods, an expectation that has been upheld in numerous examples.^{32,33,34,35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50,51,52,53,54,55,56,57,58,59,60,61,62,63,64,65,66,67,68,69,70,71,72,73,74,75,76,77,78,79,80,81,82,83,84,85,86,87,88,89,90,91,92,93,94,95,96,97,98,99,100}

This paper demonstrates the ability of an ensemble approach called Bayesian Model Averaging (BMA)⁷ to estimate solvation free energies for 45 small molecules by combining predictions from 17 diverse methods provided by the SAMPL4 challenge dataset.⁷ Though

BMA has been applied successfully for prediction tasks in many other domains^{????} this is the first application of the BMA approach to solvation energy prediction.

2 Methods

★★ check with coauthors on this paragraph ★★

This work provides a framework for addressing the uncertainties associated with estimating small molecule hydration free energies. At a high level, the framework begins by evaluating *all methods* collectively as an ensemble and uses the ensemble to make an aggregate estimate. The framework then iteratively prunes redundant or less-accurate methods from the ensemble to identify the best possible combination of methods to use for estimating hydration free energy. While improved estimates are certainly an important objective for this framework, we are equally interested in providing a better interpretation of the statistical “black box” behind these aggregate estimates. Specifically, if an aggregate estimate proves to be more accurate than individual methods, it is imperative to understand which methods are important to combine and how they are combined so that domain knowledge may be obtained from the statistical aggregation process. For example, is the optimal ensemble composed solely of implicit or explicit methods, or is it composed of a variety of different methods? If the ensemble is largely of one type of method, what are the subtle difference in the implementation of these methods (e.g., variations in small molecule force fields)? How are these methods weighted to form the aggregate? Our eventual objective with this framework is to use this type of information as a feedback loop that helps researchers with an iterative refinement of methods and their implementation.

Throughout this paper, we use the term “methods” to describe the individual approaches to calculate solvation energies and the term “models” to describe statistical ensembles of method predictions. We also define the term optimality, or optimal performance, as a forecasting method that provides the best *overall* performance based on a variety of predictive

targets: e.g., targets randomly selected through cross-validation experiments. The forecasting approach that provides the best overall performance based on mean cross-validation results is optimal in comparison to alternate approaches.

2.1 Model specification with Bayesian Model Averaging (BMA)

For estimates of hydration free energies, a basic BMA approach is to consider a set of solvation methods as a linear system.^{???} Let y_i for $i = 1, \dots, N$ be a series of hydration free energy observations for a collection of molecules, and let x_{ij} denote the i^{th} estimate obtained from the j^{th} prediction method for these observations. For example, given that y_i is the experimentally measured hydration free energy of benzaldehyde, each x_{ij} for $j = 1, \dots, P$ would be a specific method’s estimate for this value. Given P solvation methods, the combination of all x_{ij} forms the numerical ensemble estimate matrix that, along with y_i , defines a linear regression model

$$y_i = \sum_{j=1}^P x_{ij} \beta_j + \epsilon_i \tag{1}$$

Here, the parameter vector β_j defines the unknown relationship between the ensemble’s P constituents and ϵ_i is the disturbance term that captures all factors that influence the dependent variable y_i *other* than the regressors x_{ij} . Such factors include latent variables impacting the response y_i , nonlinearity effects, and measurement errors associated with quantifying y_i , as well as stochastic, experimental effects that impact the ability to reproduce y_i . Explicitly solving for the disturbance term is not feasible: ϵ_i is a theoretical, non-observable random term that must be estimated through an analysis of model residuals. Fortunately, if the distribution of $\epsilon_i : i = 1, \dots, N$ is normal (i.e., $\epsilon \sim \mathcal{N}(0, \sigma)$) the total disturbance term’s contribution to the model’s estimate goes to zero as the expectation of random variables with such distributions is $E[\epsilon|X] = 0$.

Under this assumption of normality, the objective in evaluating Equation 1 is to estimate

the values β_j that will both fit the known hydration free energy data in y_i and facilitate the ability to make inferences on the hydration free energy of unknown molecules. Many different regression techniques can estimate β_j ; however, these techniques commonly generate estimates that vary in their ability to model and infer.

BMA addresses the challenge of statistical model uncertainty by first evaluating all $2^P - 1$ possible statistical models that can be formed from the P estimation methods and combining each model’s β_j through a weighted average into an aggregated parameter vector, β_j^{BMA} (Equation 2). There are $k = 1, \dots, 2^P - 1$ distinct combinations of the P estimation methods, each with a corresponding statistical model, $M^{(k)}$, and parameter vector, $\beta_j^{(k)}$. BMA combines each $\beta_j^{(k)}$ through an average that weights each $\beta_j^{(k)}$ by the probability that its statistical model, $M^{(k)}$, is the “true” model:

$$\beta_j^{\text{BMA}} = E[\beta_j | \mathbf{y}] = \sum_{k=1}^{2^P-1} E[\beta_j^{(k)} | \mathbf{y}, M^{(k)}] \Pr(M^{(k)} | \mathbf{y}) \quad (2)$$

where $E[\beta_j^{(k)} | \mathbf{y}, M^{(k)}]$ is the expected value of the posterior distribution of $\beta_j^{(k)}$. This distribution is weighted by the posterior probability $\Pr(M^{(k)} | \mathbf{y})$ that $M^{(k)}$ is the *true* statistical model given the data \mathbf{y} . The expected posterior distribution of $\beta_j^{(k)}$ is approximated through the linear least squares solution of the given model $M^{(k)}$ and solvation energy response variable \mathbf{y} . The posterior probability term is estimated from information criteria[?]

$$\Pr(M^{(k)} | \mathbf{y}) \propto \frac{e^{-\frac{1}{2}B^{(k)}}}{\sum_{l=1}^{2^P-1} e^{-\frac{1}{2}B^{(l)}}} \quad (3)$$

where $B^{(k)}$ is the Bayesian Information Criteria for model $M^{(k)}$ estimated as

$$B^{(k)} \approx N \log(1 - R^{2(k)}) + p^{(k)} \log N. \quad (4)$$

Here $R^{2(k)}$ is the adjusted R^2 for model $M^{(k)}$ that indicates the model’s goodness of fit for the observations, $p^{(k)}$ is the number of methods used by the model (not including the intercept),

and N is the number of solvation free energy values to be predicted (i.e., the number of molecules).

The parameter vector β_j^{BMA} obtained from Equation 2 helps to address model uncertainty by accounting for all systems of linear equations that can model the relationship between the measured solvation free energy values y_i and values x_{ij} predicted by each solvation method j . Over-fitting by more complicated models is addressed through the Bayes information criterion, which penalizes models with more variables. Perhaps more importantly, β_j^{BMA} can be used to estimate new solvation free energy values for unmeasured molecules by combining new x_{ij} estimates.

2.2 Pruning models with Occam’s window

The inclusion of all $2^P - 1$ models in Equation 2 is not necessarily beneficial for predictive performance. While some models provide accurate information that will boost the ensemble’s predictive accuracy and robustness, many models will be misspecified. The cumulative effect of these misspecified models, despite the fact that they are down-weighted via low posterior probabilities, can erode the ensemble’s overall performance. Madigan and Raftery[?] developed the Occam’s Window ensemble-pruning approach that eliminates under-performing models based on Bayesian information criteria. Occam’s Window defines a set of models

$$\mathbf{A} := \{M^{(k)} \in \mathbf{M} | BIC^{(k)} - BIC^{(min)} < 6\}, \quad (5)$$

where $BIC^{(min)}$ denotes the BIC of the model $M^{(k)}$ with the lowest BIC ; low BIC indicates higher information content. The value 6 is based on Jeffreys’[?] and Raftery’s[?] assessment of Bayes factors for model comparison and ensures all $m_k \in \mathbf{A}$ meet a minimum statistical information criteria for the aggregation process. Constraining Equation 5 to the set \mathbf{A}

accelerates the evaluation of Equation 5 and improves BMA’s predictive capability: ? ?

$$\beta_j^{BMA} = E[\beta_j | \mathbf{y}] = \sum_{M^{(k)} \in \mathbf{A}} E[\beta_j^{(k)} | \mathbf{y}, M^{(k)}] Pr(M^{(k)} | \mathbf{y}) \quad (6)$$

BMA estimates method j ’s utility for explaining a set of observations, \mathbf{y} , by assessing the probability that the method’s coefficient term, β_j^{BMA} will receive a non-zero value. The estimate of this probability is based on the conditioned, cumulative sum of all model posteriors

$$Pr(\beta_j^{BMA} \neq 0) = \sum_{M^{(k)} \in \mathbf{A}} Pr(M^{(k)} | \mathbf{y}) I(j) \quad (7)$$

where

$$I(j) := \begin{cases} 1, & \text{if model } M^{(k)} \text{ specifies } j \text{ as a regressor} \\ 0, & \text{otherwise.} \end{cases} \quad (8)$$

Equation 7 can be used to sort and prioritize the methods based on their utility; i.e., the probability that the coefficient term weighting a method’s estimate will not be 0. The combination of Equations 6 and 7 therefore provide a statistical framework to support an iterative, statistical process for designing an ensemble that will provide the best combination of models and methods for estimating solvation free energy.

2.3 Statistical design of method ensembles

At a summary level the statistical design process, shown in Algorithms 1 and 2, proceeds in two iterative stages. The first stage, presented in Algorithm 2, is the analysis stage and its function is to evaluate the statistical significance of each method in the current version of the ensemble via Equation 7. This information is based on a set of 2-fold cross-validation tests that randomly create training and validation datasets from the small molecule compounds in the SAMPL4 challenge. In these tests, each method’s statistical information is thus evaluated against a variety of random training and validation sets. The results of this first stage are

Algorithm 1**Prune and Design Ensemble**

Require: vector $\mathbf{y}^{n \times 1}$, matrix $\mathbf{x}^{n \times p}$

```

1: loopCount = 1
2: rmse = null
3: x_solution =  $\mathbf{x}$ 
4:  $\beta_{BMA}^{1 \times j} = \mathbf{0}^{1 \times j}$ 
5: while (loopCount < p) do
6:    $\beta_{BMA}, s, \mathbf{v} = \text{AssessMethods}(\mathbf{y}, \mathbf{x})$ 
7:   if (s < rmse) or (rmse == null) then
8:     x_solution =  $\mathbf{x}$ 
9:     rmse = s
10:  end if
11:  m = 1
12:  for (methods  $\in \mathbf{x}$ ) do
13:    if ( $\mathbf{v}[\text{methods}] < \mathbf{v}[m]$ ) then
14:      m = methods
15:    end if
16:  end for
17:   $\mathbf{x} = \mathbf{x}.\text{delete}(m)$  {prune least significant method}
18:  loopCount += 1
19: end while
20: return x_solution,  $\beta_{BMA}$ , rmse

```

Algorithm 2**Analyze Methods**

Require: vector $\mathbf{y}^{n \times 1}$, matrix $\mathbf{m}^{n \times j} : j \leq p$

```

1: cv = 100 {100 cross-validation experiments}
2:  $\beta_{BMA}^{1 \times j} = \mathbf{0}^{1 \times j}$ 
3: error = 0
4:  $\mathbf{v}^{1 \times j} = \mathbf{0}^{1 \times j}$ 
5: while (index < cv) do
6:    $\mathbf{m}_t^{q \times j}, \mathbf{m}_v^{r \times j} = \mathbf{m}^{n \times j}$  {split  $\mathbf{m}$  into train and
    validate}
7:    $\beta_{BMA} = \beta_{BMA} + \text{BMA}(\mathbf{m}_t)$  {Equation 6}
8:   error = error +  $\text{RMSE}(\mathbf{Y}, \mathbf{m}_v \times \beta_{BMA})$ 
9:    $\mathbf{v} = \mathbf{v} + \text{P} \neq 0(\mathbf{m}_t)$  {Equation 7}
10:  index = index + 1
11: end while
12: return  $\frac{\beta_{BMA}}{cv}, \frac{\text{error}}{cv}, \frac{\mathbf{v}}{cv}$ 

```

then sent to stage two (Algorithm 1): the ensemble pruning stage.

The second stage evaluates the performance of the current ensemble and then prunes this ensemble by removing that method with the least statistical significance. The process then repeats using the new, pruned ensemble. By monitoring the performance of each ensemble across the pruning process, the ideal set of methods is able to be identified.

At a more detailed level, the design process assumes an initial set of n estimates made by p methods, $\mathbf{x}^{n \times p}$ for observations, $\mathbf{y}^{n \times 1}$. The second and third lines initialize the root mean squared error (RMSE) and the solution for the best ensemble of models and methods. The algorithm will iterate over the following tasks while the algorithm still has more than two methods to evaluate (line 5 - line 19).

First, the performance of the current ensemble is evaluated through *AssessMethods* (line 6). The *AssessMethods* function, shown in Algorithm 2, takes as input the same observation data, \mathbf{y} , as Algorithm 1, as well as a set of N estimates made by a subset of $j \leq p$ methods. Algorithm 2 performs 100 iterations of a 2-fold cross-validation and begins by partitioning the estimate matrix \mathbf{m} equally into training and validation data (line 6). Using the training

data, the algorithm uses Equation 6 to estimate β_{BMA} and then calculates the RMSE of this coefficient vector based on the validation data (lines 7 and 8). Next, the $\Pr(\beta_j^{BMA} \neq 0)$ of each method are calculated and saved in the vector, \mathbf{v} (line 9). After 100 iterations, the function returns an estimate for β_j^{BMA} , the mean RMSE of the current ensemble, as well as a list of all methods and their corresponding probability values for $\Pr(\beta_j^{BMA} \neq 0)$. Next, Algorithm 1 compares and conditionally updates the current “best” performing model, **x_solution**, with the model returned by *AssessMethods* (lines 8 and 9). The algorithm then identifies the method with the lowest $\Pr(\beta_j^{BMA} \neq 0)$ and removes this method from the current model (line 17). With this method removed, the process repeats until there are just two methods left. The output of this process is an ensemble of methods, \mathbf{x}_{ij} that are statistically determined to be the best methods for estimating the observations in \mathbf{y} and a statistical model, β_{BMA} that specifies how to best combine these methods.

2.4 Solvation free energy data and solvation methods

The SAMPL4 challenge consists of 49 submissions representing a total of 19 different research groups.[?] Each of these methods provides solvation free energy estimates for 45 small molecule compounds.¹ The challenge is blind in the sense that the solvation free energy values for these molecules were hidden from participants; e.g., energy values are not found in standard solvation free energy test sets, and their values are not readily available in the literature.[?] In this work, we restrict our analysis to a subset of 17 method submissions based on the fact that many groups made multiple submissions that were strongly correlated. In these cases, we chose only a single variant to ensure that multicollinearity did not inflate the significance of specific methods during model selection and averaging; such bias can create unstable estimates for β_j^{BMA} and that in turn can reduce BMA’s estimate accuracy.[?] These methods are summarized in Table I based on the methods used to calculate the solvation

¹The original SAMPL4 study used 52 samples; thus some samples have labels like SAMPL4_052. However, several compounds were removed during the course of the study based on “problems with experimental values, SMILES strings, or structures for a number of compounds, resulting in removal of some compounds from the challenge.”[?]

Table I: This table lists the solvation methods used in our ensemble design process. Method ID indicates the identification number of the method that is referenced throughout this paper. Sampling strategies include quantum mechanical (QM), molecular dynamics (MD), and molecular mechanics with Poisson-Boltzmann surface area solvation (MM-PBSA). The listed performance is based on 100 iterations of a 2-fold cross-validation study as described in the text. This performance is also shown graphically in Figure 1. The last column is a comparison of each method to the optimal BMA ensemble, BMA (Stage 16). This column indicates that the ensemble design approach presented in this work is able to reduce estimation errors by 29% to 91% (i.e., 1.2–9.4 kcal mol⁻¹) in comparison to the individual methods. The final BMA ensemble is indicated by the two italicized methods: *imp-2* and *alc-3*. Wilcoxon based *p*-values for BMA’s mean RMSE distribution vs. the best method’s mean RMSE distribution (*imp-2*) are shown in Table II.

Reference	Method ID	Methodology	Ensemble Mean RMSE (kcal mol ⁻¹) and Std Err Mean	Improvement with BMA (Stage 16)
Coleman et al. [?]	imp-6	multi-conformation implicit	9.40 ± 0.2	91%
Parsod et al.	alc-5	alchemical MD	4.50 ± 0.071	82%
Sharp et al. [?]	imp-4	single conformation implicit	3.80 ± 0.045	78%
Jiafu et al.	exp-3	QM/MD	2.80 ± 0.045	71%
Purisima et al. [?]	imp-5	single conformation implicit	2.55 ± 0.049	68%
Mark et al. 2011	alc-2	alchemical MD	2.40 ± 0.036	66%
Genheden et al. [?]	exp-1	MM-PB/SA	2.00 ± 0.022	59%
Weyang et al.	exp-4	MM-PB/SA	1.84 ± 0.031	55%
Biorga et al. ^{??}	alc-4	alchemical MD	1.65 ± 0.013	50%
Elingson et al. [?]	imp-7	single conformation implicit	1.52 ± 0.039	46%
Fennell et al. [?]	hyb-2	hybrid	1.52 ± 0.013	46%
Jambeck et al. [?]	alc-1	alchemical MD	1.52 ± 0.020	46%
Park [?]	imp-3	single conformation implicit	1.44 ± 0.020	43%
Klamt et al. [?]	imp-1	single conformation implicit	1.36 ± 0.030	40%
<i>Gilson et al.[?]</i>	<i>alc-3</i>	<i>alchemical MD</i>	1.24 ± 0.017	34%
Geballe et al. [?]	imp-8	single conformation implicit	1.17 ± 0.017	30%
<i>Sandberg et al.[?]</i>	<i>imp-2</i>	<i>multi-conformation implicit</i>	<i>1.15 ± 0.023</i>	<i>29%</i>
BMA (Stage 16)	NA	ensemble	0.82 ± 0.015	0

energy:

- **Group 1:** Single-conformation implicit solvent methods^{???} based on Poisson-Boltzmann and related methods;^{???}
- **Group 2:** Multi-conformational implicit solvent methods;^{?????}
- **Group 3:** Methods based on molecular dynamics-based free energy calculations in explicit solvent^{????} with small molecule force fields;[?] and,
- **Group 4:** Hybrid solvent methods.[?]




Figure 1: This figure depicts the mean root mean squared error (RMSE) for the 17 initial methods used in our ensemble design process. Additionally, the 95 % confidence interval (based on standard error) for the expected mean RMSE is shown for each method. Thus at 95% statistical confidence, the true expected performance for each method is depicted according to the 100 iterations of the 2-fold cross-validation experiment we detail in Algorithm 2.

2.5 Training, estimating, and pruning an ensemble with BMA

From this initial set of 17 methods, we iteratively assessed and pruned methods as described in Algorithm 1. The approach began by randomly sampling (without replacement) 26 of the original 52 experimental solvation free energy measurements to form training and validation datasets (Algorithm 2, line 6). Collectively, these sampled values formed the observation vector y_i and the estimates from each of the prediction methods in Table I for these measurements defined the ensemble estimate matrix, x_{ij} . The observation vector and the ensemble estimate matrix were used to form the linear system in Equation 1.

Next, we estimated the β_j^{BMA} parameter from Equation 2 by assessing all possible $(2^{17}-1)$ statistical models $M^{(k)}$. Each model’s information criteria $B^{(k)}$ was used to identify a reduced set of the most informative models per Equation 5. Based on this reduced modeling space, the coefficient terms β_j^{BMA} were estimated through weighted averages of each statistical model’s ordinary least squares solution (Equation 6). We used β_j^{BMA} to estimate the remaining 26 solvation free energy measurements that were *not* used to train the BMA model. This task was accomplished by combining the estimates of all methods in Table I for the validation data with β_j^{BMA} to produce an aggregated estimate. An RMSE was obtained for the 26 estimates made on validation data.

This process was repeated 100 times (Algorithm 2) so that performance for any estimation method could be reported as a mean RMSE: i.e., the mean of 100 RMSEs that each represent performance for a 2-fold cross-validation using the 26 member validation set described above. In addition to the mean RMSE, the information content provided by each method in the ensemble is also returned.

Algorithm 1 used this information to perform two tasks. First, if the performance of this aggregated estimate was better than all previously examined aggregates, the statistical model combining these methods was saved as the new optimal model (Algorithm 1, lines 7-9). Next, Algorithm 1 pruned a method from the ensemble that provided the least amount of information to the aggregated estimate. The pruning process was repeated until just two methods are left (Algorithm 1, lines 12-18). The ensemble of methods whose aggregate forecast has the lowest mean RMSE was saved and returned as the final model for the ensemble. This final model, and the set of methods that correspond to this model, were the final products of the statistically driven ensemble design process, referred to as the BMA-based optimal ensemble.

3 Results and Discussion

We report the results of our BMA-based ensemble in three stages. First, we contrast the performance of the initial 17 SAMPL4 methods to the performance of the BMA-based optimal ensemble. We then present data on the iterative design process used to create this optimal ensemble. Second, we examine the *conditional* performance of this optimal ensemble according to the individual molecules used in the SAMPL4 challenge. In this stage, we show how the BMA-based optimal ensemble provides more reliable estimates in comparison to individual methods, especially against the more challenging set of small molecules from the SAMPL4 dataset. Finally, we complete the analysis of the optimal ensemble by contrasting its performance to the performance of alternate statistical techniques that can be used to form aggregated estimates.

In all analyses, we report the performance based on the mean RMSE associated with the 100 X 2-fold cross-validation. Further, we report the standard error of the mean (SEM) to better characterize the expected performance for these estimates. In contrast to standard deviation, which quantifies the spread of the cross-validation distribution, the SEM metric directly maps cross-validation performance into confidence intervals that bound the true expected performance of each RMSE.

We report the statistical significance of all performance data through a Wilcoxon rank sum paired comparison test.[?] This non-parametric approach tests the hypothesis that the mean RMSE distributions of two approaches are equal: $H_0 : \mu_Y = \mu_X$. To control the family-wise error rate of our tests, we applied a Bonferroni correction to determine corrected p -values with a threshold of $\alpha = 0.05$. Thus when comparing BMA to a given method, a Wilcoxon generated p -value greater than 0.05 indicates we fail to reject H_0 : the distributions are thus equal and we conclude that BMA and the method are equivalent in their performance. On the other hand, Wilcoxon generated p -values that are less than 0.05 indicate we should reject H_0 . In this latter case, we then compare the mean RMSE for BMA and the given method to assess performance.

Finally, a Shapiro-Wilk test[?] was used to test the hypothesis, H_0 , that residuals in the SAMPL4 challenge are not significantly different from a normal population. Under an $\alpha = 0.05$, a p-value statistic of 0.9556 was obtained from this test, indicating that there was insufficient information to reject H_0 . Without evidence to reject H_0 , assumptions of normality indicate that the expectation for the disturbance terms ϵ_i in Equation 1 (e.g., sources associated with measurement error, stochastic effects, etc.) sum to 0 and are ignored.

3.1 Comparing estimates from BMA’s optimal ensemble to SAMPL-4 challenge methods

The performance of all methods used in this work is shown in Table I and Figure 1. Column 3 in Table I lists the specific method for each estimation approach. The performance results in column 4 of this table are consistent with results reported by Mobley et al.[?] The performance of the optimal ensemble is shown in the last row. The corresponding methods that are constituents in this optimal ensemble, “alc-3” and “imp-2”, are highlighted in blue. The final column in Table I lists a direct comparison of each method’s performance to the performance of the optimal ensemble: e.g., the ensemble reduces estimation errors by as much as 91% in comparison to imp-6 and by 29% in comparison to imp-2.

The box plots in Figure 1 illustrate the performance variability across the different methods. These plots also underscore the amount of uncertainty inherent in estimating solvation free energies: method-selection uncertainty plays a confounding factor in estimating solvation free energies. Finally, the red line indicates the mean RMSE of the best-performing method: imp-2. This line is also used in Figure 2 to contrast the iterative improvements obtained during the optimal ensemble’s design process. The ensemble’s iterative design process based on Algorithms 1 and 2 is shown in Table II and Figure 2. The process starts with the 17 initial methods shown in Table I. Each subsequent row in Table II represents an iteration through the pruning process. The second column lists the mean RMSE of each stage’s ensemble based on its aggregated estimate. At the end of each stage, a method is




Figure 2: This plot illustrates the iterative pruning process discussed in Section 2.2. The y-axis depicts the mean root mean squared error (RMSE) of the different aggregated estimates based on ensembles formed during the different iterations of pruning. The x-axis indicates the stages of pruning. In general, the mean RMSE and the 95% confidence intervals for each mean (based on standard error for the mean SEM) reduce with each iteration. The performance of the different ensembles are compared to the best performing method through the red line at $y = 1.15$; all iterations past Stage 7 outperform the best method in the ensemble. Based on Wilcoxon generated p-values, the significance in the distributions of mean RMSE between the different ensembles are presented in Table II. Based on mean RMSE, the optimal ensemble is created at Stage 16.

selected to be pruned from the existing ensemble before proceeding to the next iteration; the specific method that was selected is shown in column 3. For example, at Stage 1 there were 17 methods in the ensemble and imp-6 was selected to be removed. During the next step, Stage 2, imp-6 was removed from the ensemble so that there were only 16 methods used to create an aggregated estimate. At the end of this stage, alc-4 was selected to be pruned for Stage 3. In the final stage, the only remaining methods in the ensemble were imp-2 and

Table II: This table lists the performance of the aggregated estimates obtained from the different ensembles created during the design process in Section 2.2. The second column lists the mean root mean squared error (RMSE) for each ensemble’s aggregated estimate based on 100 iterations of a 2-fold cross-validation; this performance is also shown in Figure 2. The third column lists the method that was selected to be pruned from the ensemble at the *next* stage of the design process. The Wilcoxon generated p -values in column four are based on comparisons of mean RMSE distributions obtained from sequential ensembles and their aggregated estimates. p -values that are greater than 0.05 are shown and indicate distributions that are equivalent (e.g., Stage 3 and 2) based on an $\alpha = 0.05$. Similarly, column five lists the Wilcoxon generated p -values reflecting comparisons between each stage and the best performing method, imp-2. As with column 4, bold p -values indicate that the performance between imp-2 and a given stage are equivalent (e.g., Stages 4-6). Based on mean RMSE and p -values from this table, the optimal ensemble is the one created in Stage 16; the final column in this table lists the performance improvement this ensemble provides in comparison to the ensembles generated in previous stages.

Stage	Ensemble RMSE (kcal mol ⁻¹) Mean and Std Err Mean	Method Selected for Pruning	p -value (sequential stages)	p -value (stage vs. imp-2)	Improvement with BMA (Stage 16)
1	1.40 \pm 0.060	imp-6	NA	0.000	41%
2	1.32 \pm 0.053	alc-4	0.411 (vs. Stage 1)	0.002	38%
3	1.29 \pm 0.046	imp-8	1.000 (vs. Stage 2)	0.004	36%
4	1.19 \pm 0.031	imp-5	0.012 (vs. Stage 3)	0.356	31%
5	1.14 \pm 0.029	imp-3	0.000 (vs. Stage 4)	0.438	28%
6	1.13 \pm 0.026	imp-7	0.883 (vs. Stage 5)	0.306	27%
7	1.09 \pm 0.025	alc-5	0.015 (vs. Stage 6)	0.027	25%
8	1.06 \pm 0.019	alc-2	0.055 (vs. Stage 7)	0.003	23%
9	1.03 \pm 0.020	imp-1	1.000 (vs. Stage 8)	0.001	20%
10	0.99 \pm 0.017	imp-4	0.000 (vs. Stage 9)	0.001	17%
11	0.96 \pm 0.018	alc-1	0.000 (vs. Stage 10)	0.000	15%
12	0.93 \pm 0.018	exp-4	1.000 (vs. Stage 11)	0.000	12%
13	0.90 \pm 0.017	hyb-2	1.000 (vs. Stage 12)	0.000	09%
14	0.90 \pm 0.018	exp-1	1.000 (vs. Stage 13)	0.000	09%
15	0.91 \pm 0.018	exp-3	0.001 (vs. Stage 14)	0.000	10%
16	0.82 \pm 0.015	NA	0.000 (vs. Stage 15)	0.000	0%

alc-3. This iterative design process is also graphically illustrated in Figure 2 showing how selective pruning increases the performance of each successive ensemble. The red line in this figure indicates the performance of the best performing method, imp-2. The benefits of the aggregated estimates become apparent after Stage 6 where the ensembles outperform imp-2.

The statistical significance of the iterative design process is listed in columns 4 and 5 in Table II; bold p -values in these columns indicate the performance between two distributions are equivalent. Column 4 lists the Wilcoxon generated p -values based on comparisons of mean RMSE distributions obtained from sequential ensembles. For example, in the second

row the p -value for the comparison of Stage 2 vs. Stage 1 indicates that these distributions are equivalent. Contrariwise, in row 4, the p -value for the comparison between Stage 4 and Stage 3 indicates that the distributions are not equivalent; the mean RMSE in column 2 indicates that the ensemble of Stage 4 outperforms the ensemble built in Stage 3.

As a second analysis of significance, column 5 lists the Wilcoxon generated p -values that represent comparisons between each ensemble and the best-performing method, imp-2. From these values, Stages 4-6 are seen to be equivalent to imp-2. However, the mean RMSE distributions of all subsequent stages are not equivalent: based on mean RMSE listed in column 2, we conclude that these successive ensembles (increasingly) outperform imp-2. The p -values in these columns and the mean RMSE results demonstrate that the optimal ensemble is created in Stage 16. The final column in Table II lists the performance improvement that the Stage 16 ensemble provides in comparison to all the other ensembles generated in the design process.

Figure 3 depicts a heat-map that shows the statistical information that drove the pruning process for each stage. In this image, the y-axis represents the different methods; the x-axis (starting from left) indicates the successive stages in the design process. The color scale represents the mean probability, 0 - 100%, that a given method’s coefficient term, β_j will not be zero. At Stage 1, all methods were used in the ensemble and their $\Pr(\beta_j^{\text{BMA}} \neq 0)$ ranged from 40% (imp-6) to 80% (imp-2). As imp-6 had the lowest mean probability of not being zero, imp-6 was pruned after Stage 1; the color map colors this method white in Stage 2 to indicate it has been eliminated. The methods listed on the legend of the y-axis are ordered by the sequence that they were eliminated in the design process: imp-6 first, alc-4 second, imp-8 third, etc.

3.2 Performance analysis based on compounds

Figures 4 and 5 depict the performance of different methods according to specific SAMPL4 challenge small molecule compounds. In addition to the optimal ensemble, we also show




Figure 3: This image provides a graphical depiction of the ensemble design process performed on the initial ensemble of 17 methods. The color scale represents the probability, 0 - 100%, that a given method’s coefficient term, β_j will not be zero. The stages in the design process, starting at the left and moving to the right, are shown on the x-axis. The methods in each ensemble are shown on the y-axis; note that the methods are listed (from top to bottom) in the order that they are pruned in the design process. Thus at Stage 1, all methods are used in the ensemble and their $\Pr(\beta_j^{\text{BMA}} \neq 0)$ range from 40% (e.g., imp-6) to 90% (imp-2). By Stage 3, imp-6 and alc-4 have been pruned from the ensemble and the probability values have adjusted accordingly as shown in the colored column above Stage 3. In general the trend across the different stages of pruning illustrates that methods become increasingly lighter, i.e., the ensemble design process becomes increasingly confident that these methods (e.g., imp-7 and exp-4) are not needed in the ensemble. Contrariwise, the $\Pr(\beta_j^{\text{BMA}} \neq 0)$ for a few methods (e.g., alc-3, and imp-2) remain above 70% and even increase throughout the design process, indicating a high degree of confidence in the statistical significance of these methods.

the performance of the first, second, and third best-performing methods from the SAMPL4 challenge: imp-2, imp-8, and alc-3. Methods imp-2 and alc-3 are the methods used in the optimal BMA ensemble and exp-3 is the final method eliminated from the ensemble (Stage 15 in Table II).

Mobley et al.[?] note that predictive error generally increases as molecular complexity increases. Specifically, polyfunctional molecules with several interacting groups were especially difficult to estimate in the SAMPL4 challenge: SAMPL4_022 (mefenamic acid),




Figure 4: This figure is one of two figures (Figures 4 and 5) that depict the RMSE (kcal mol^{-1}) performance of several methods based on the individual compounds taken from the SAMPL4 challenge: the first, second, and third best-performing methods (i.e., imp-2, imp-8, and alc-3) as well as exp-3. Note that imp-2 and alc-3 are the methods used in the optimal BMA ensemble (Stage 15) and exp-3 is the final method eliminated from the ensemble (Table II) in Stage 16. BMA’s performance based on the optimal ensemble is shown based on its distribution of the mean RMSE for estimates made in our 2-fold cross-validation analysis. Of note is the performance of the different methods for SAMPL4.022 (mefenamic acid), SAMPL4.023 (diphenhydramine), SAMPL4.027 (1,3-bis-(nitroxy)propane), SAMPL4.009 (2,6-dichlorosyringaldehyde), and SAMPL4.001 (mannitol). These are the most challenging compounds for methods to estimate based on the analysis of Mobley et al. of the SAMPL4 data.[?] The benefits of the ensemble is clearly demonstrated here as the BMA ensemble outperforms all methods in estimating SAMPL4.022, SAMPL4.009, and SAMPL4.001. For SAMPL4.023 and SAMPL4.027 the ensemble provides the second best performance, and in this context provides more consistent performance than the other methods: e.g., alc-3 is better at estimating SAMPL4.027, but is third at estimating SAMPL4.23.




Figure 5: This figure is one of two figures (Figures 4 and 5) that depict the RMSE (kcal mol⁻¹) performance of several methods based on the individual compounds taken from the SAMPL4 challenge: the first, second, and third best-performing methods (i.e., imp-2, imp-8, and alc-3) as well as exp-3. See Figure 4 for more information about both plots.

SAMPL4.023 (diphenhydramine), SAMPL4.027 (1,3-bis-(nitroxy)propane), SAMPL4.009 (2,6-dichlorosyringaldehyde), and SAMPL4.001 (mannitol). Figures 4 and 5 show that the optimal ensemble outperforms all methods in estimating SAMPL4.022 and SAMPL4.001, and provides the second best estimates for SAMPL4.009, SAMPL4.023 and SAMPL4.027. Aside from the optimal ensemble, there is no clear best method for estimating these compounds. For example, while imp-8 is best at estimating SAMPL4.009, it does not do well at estimating either SAMPL4.023 or SAMPL4.027. Similarly while imp-2 performs well at estimating SAMPL4.023, it does not do as well at estimating SAMPL4.009 or SAMPL4.027. The general trend in performance indicates that the optimal ensemble provides more consistent and accurate estimates than any specific method.

3.3 Interpreting aggregate estimate performance

Aggregate estimate performance is degraded by both the predictive sensitivities of its method constituents and the uncertainties associated with statistical model construction. These sources of error manifest in different ways.

If methods are sensitive to functional discrepancies between training and testing molecules, then aggregate estimates that rely on these methods may perform poorly on chemical compounds that are outside the training set. While some of these sensitivities may be obvious to identify (e.g., methods that are only accurate for certain chemical classes), most are much more subtle and difficult to detect. Cross-validation tests, like those performed in Section 2.3, help assess this first, more obvious case where there are broad sensitivities based on class or functional groups. More specifically, under cross-validation tests the overall performance of the aggregate estimate will be significantly poorer in comparison to other methods. This trait will be most notable when comparing the aggregate against its method constituents on molecules where constituents are very accurate. Additionally the 95% confidence intervals of the expected mean for most chemical compounds will be much wider, comparatively, to the most narrow confidence intervals.

In this work the overall performance of the aggregate estimate is overall better than its constituents (Table 1, Figures 4 and 5,) and significantly better than other methods. Additionally, the 95% confidence intervals for the expected mean for the majority of the chemical compounds are very narrow (especially in comparison to SAMPL4_001). Based on these observations, large method-based sensitivities do not appear to be an issue for the SAMPL4 challenge data and so we next proceed to assess the sensitivities associated with statistical model construction.

The design of a statistical ensemble relies on the assumption that some combination of known methods exist that will form an accurate aggregate estimate across all models in the cross-validation study. A poorly designed model

design and tuning process attempts to construct a statistical model that has learned the predictive strengths of its constituents, and how to combine these strengths to form an aggregated estimate. Success in combining these methods to make accurate estimates relies on the assumption that some combination of methods does exist and if the model for combining these methods was known the aggregate estimate would be

A perfect ensemble of methods is a statistical model that has learned: 1) the predictive performance trends of its methods; and, 2) how to combine these trends to form a perfect estimate for any chemical compound. The gold standard of a perfect ensemble is not a realistic expectation, however, an optimal ensemble can provide some benefits towards this goal. Additionally, by iteratively assessing performance and refining the ensemble and its methods, marked increases in estimate accuracy can occur over time.

The second effect is that the 95% confidence interval for the expected mean (or, alternately standard deviation) will be much wider, comparatively for compounds where there are method sensitivities than for

expected aggregate estimate’s performance when estimating any compound in a dataset, *regardless* of which molecules were used to train and which were used to test. In this sense, cross-validation approaches help to address some of the sensitivities that may result from

distinctly different training and testing datasets.

Methods that accurately predict certain classes of compounds, but generally predict poorly on other classes have an inconsistency that can effect aggregate performance in two ways. First, ensembles designed on training data that overemphasize the predictive performance of certain methods may potentially is largely representative of

and we broadly characterize four types of behavior based on .

- Type 1:
 - **Observation:** The aggregated estimate bests at least one or more constituent. 95% confidence intervals of the expected mean are **narrow**.
 - **Interpretation:** The ensemble’s underlying model is valid. Method predictions are consistent with other Type 1 molecular targets.
 - **Examples** (Figures 4 and 5): SAMPL4_003 - SAMPL4.016
- Type 2:
 - **Observation:** The aggregated estimate bests at least one or more constituent. 95% confidence intervals of the expected mean are **wide**.
 - **Interpretation:** Uncertainty with underlying model based on training/testing sensitivity. Method predictions are consistent with other Type 1 molecular targets.
 - **Examples** (Figures 4 and 5): SAMPL4_001 and SAMPL4_045
- Type 3:
 - **Observation:** The aggregated estimate does not beat any constituent. 95% confidence intervals of the expected mean are **narrow**.
 - **Interpretation:** There is strong consistency in method predictions, however the ensemble’s underlying model is challenged.

- **Examples** (Figures 4 and 5): SAMPL4_047
- Type 4:
 - **Observation:** The aggregated estimate does not beat any constituent. 95% confidence intervals of the expected mean are **wide**.
 - **Interpretation:** There is strong consistency in method predictions, however the ensemble’s underlying model is slightly challenged
 - **Examples** (Figures 4 and 5):

A perfect ensemble of methods is a statistical model that has learned: 1) the predictive performance trends of its methods; and, 2) how to combine these trends to form a perfect estimate for any chemical compound. The gold standard of a perfect ensemble is not a realistic expectation, however, an optimal ensemble can provide some benefits towards this goal. Additionally, by iteratively assessing performance and refining the ensemble and its methods, marked increases in estimate accuracy can occur over time.

In returning to the performance seen in Figures 4 and 5, the optimal ensemble outperforms its constituents in over 40 small molecule compounds.

The remaining molecules challenge the ensemble’s estimate as both of its method constituents (alc-3 and imp-2) outperform the ensemble for SAMPL4_002, SAMPL4_017, SAMPL4_024, SAMPL4_030, and SAMPL4_047. In three of these cases, SAMPL4_002, SAMPL4_017, and SAMPL4_030, the aggregate’s estimate is only just outperformed by its constituents. This degree of error is small enough to potentially be associated with small modeling uncertainties in the estimation of the β_j terms. However, aggregate estimate performance for SAMPL4_024 and SAMPL4_047 present very different cases. In both of these cases, the aggregate estimate is clearly not accurate.

INTERPRETING THE COMPOUNDS: 47 is the worse performer for the aggregate. These two compounds are the largest MW compounds evaluated by the SAMPL4 challenge:

SAMPL4.024 is amitriptyline (MW 277.41) and SAMPL4.047 is 1-(2-hydroxyethylamino)-9,10-anthraquinone (MW 267.28). Additionally, SAMPL4.047 possesses the second largest hydration free energy in the challenge.

First, if there are large functional discrepancies between training and testing molecules then aggregate estimates can perform poorly.

The second source of error is based on a lack of consistency in method predictive performance. In general, if methods in the optimal ensemble are largely *consistent* in their predictive performance, the aggregate estimate’s performance should not be effected if there are large functional discrepancies between training and testing datasets. For example, if the methods in the optimal ensemble are all moderately strong performers

: i.e., stronger performing methods in the training sets are also the stronger performing methods in the testing sets. This work’s statistical framework and the approach presented in Algorithms 1 and 2 help to robustly evaluate and identify these methods and how to best combine them an aggregate estimate. However, note while the framework will provide the best ensemble possible based on existing methods, the framework may be able to address all method uncertainties. In these cases, the optimal ensemble may actually perform comparatively to its constituents, or even worse than all constituents.

often confused with the first source of error. is only partially correct as there is a deeper cause that is the more significant factor to assess in examining aggregate performance.

Cross-validation tests, like those performed in Section 2.3, help to assess the expected aggregate estimate’s performance when estimating any compound in a dataset, *regardless* of which molecules were used to train and which were used to test. In this sense, cross-validation approaches help to address some of the sensitivities that may result from distinctly different training and testing datasets.

If the 95% confidence intervals for this expected performance are close to the mean then concerns for a poorly performing estimate likely has another source. When assessing an ensemble’s performance, a second attribute to characterize is the degree that ensemble

constituents provide *consistent* performance over varied chemical targets.

3.4 Performance analysis of alternate ensemble techniques

There are other approaches beside BMA that can combine an ensemble of methods into an aggregate estimate. In our cross-validation study, we evaluated four common approaches for aggregating an ensemble and evaluated their estimated benefits in comparison to BMA. These methods are listed in Table III and include: Random Forest,[?] Ridge Regression,[?] Lasso,[?] and stepwise regression via forward selection. These techniques were chosen as they

Table III: This table lists the performance of different ensemble approaches in comparison to the optimal ensemble designed in this work through BMA. The performance of these ensemble approaches, given as the mean RMSE with standard deviation, are based on the 100 iterations of the 2-fold cross-validation experiment discussed in Section 2.2. Based on an $\alpha = 0.05$, the Wilcoxon based p -values indicate that BMA’s improved performance is statistically significant to the other ensemble approaches for combining methods to make an aggregated estimate. The last column indicates the improvement in estimation that the optimal ensemble provides to these alternate techniques: estimation accuracy is improved from 25% to 61%.

Ensemble Model	Ensemble Mean RMSE and Std Err Mean	p-value	Improvement by BMA (Stage 16)
Random Forest	2.08 \pm 1.00	0.00	61%
Ridge	2.06 \pm 0.75	0.00	60%
Lasso	1.12 \pm 0.34	0.00	27%
Forward Selection	1.09 \pm 0.26	0.00	25%
BMA (Stage 16)	0.82 \pm 0.17	NA	0%

have all been used successfully for a variety of inference tasks and are readily available for use.[?] ? As all of these approaches construct an estimate for β , training and estimating with these approaches was performed identically to BMA, using the cross-validation detailed in Section 2.2. We also followed the same procedure for comparing BMA’s estimating capability to these alternate ensemble-based techniques. Figure 6 and Table III provide an overview of the cross-validation errors for the various ensemble-based approaches and the BMA-based optimal ensemble. The statistical significance of BMA’s performance in Figure 6 is based on p -values shown in Table III. Based on an $\alpha = 0.05$, Table III indicates that we reject the




Figure 6: This figure displays the performance of different ensemble approaches in comparison to the optimal ensemble designed in this work, BMA (Stage 16). The mean RMSE, min, max, first, and third quartiles (kcal mol^{-1}) of these ensemble approaches are shown based on the 100 iterations of the 2-fold cross-validation experiment discussed in Section 2.2. Based on an $\alpha = 0.05$, the Wilcoxon based p-values (Table III) indicate that BMA’s improved performance is statistically significant to the other approaches that can combine an ensemble of methods to make an aggregated estimate.

null hypothesis for all paired comparison tests. BMA’s mean RMSE distribution is therefore not equivalent to the mean RMSE distribution of any other ensemble-based technique. As the distributions are not equal, we compared mean RMSE distributions of BMA to the other ensemble-based approaches in Figure 6 and Table III. From these mean RMSEs, it is clear that the BMA-based approach outperforms all other ensemble-based estimation approaches: BMA-based estimates reduced error by approximately 60% in comparison to Random Forest and Ridge regression methods. In comparison to Lasso, BMA reduces estimation error by approximately 27%. Finally, in comparison to stepwise regression via forward selection, BMA reduces error by approximately 25%.

4 Conclusions

This study demonstrates a proof-of-principle application of how to statistically design and aggregate an ensemble of methods for estimating solvation free energies in small molecules. While the performance of BMA is expected to generalize to a much broader set of small molecule estimation problems, the specific BMA model trained in this study is likely to be dependent on the small molecules used in the SAMPL4 challenge. The BMA approach is purely statistical in nature, but could also be trained to modify the aggregation process based on structural and environmental features (e.g., only look at ensembles of empirical methods for certain structural features and consider all methods for other structures). In future work we will look at penalizing computationally expensive methods that provide minimal accuracy benefits to optimize both the accuracy of the estimates and the efficiency of the calculation.

Acknowledgments

This research was funded by NIH grant R01 GM069702 to NAB.